

Employee Attrition Prediction

Student: Amruta Kumbar
Roll Number: AA.SC.P2MCA24074031
Course: MCA Minor Project (21CSA697A)
Date: February 2026

This notebook contains the complete code for Employee Attrition Prediction using Machine Learning and Explainable AI.

Table of Contents

1. [Import Libraries](#)
2. [Load Dataset](#)
3. [Data Exploration](#)
4. [Data Preprocessing](#)
5. [Feature Engineering](#)
6. [Train-Test Split](#)
7. [Model Training](#)
8. [Model Evaluation](#)
9. [SHAP Explainability](#)
10. [Visualization](#)
11. [Save Model](#)
12. [Make Predictions](#)

1. Import Libraries

```
# =====
# STEP 1: IMPORT ALL REQUIRED LIBRARIES
# =====

# Data manipulation
import pandas as pd
import numpy as np

# Visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Machine Learning
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    roc_auc_score, confusion_matrix, classification_report, roc_curve
)

# XGBoost
import xgboost as xgb

# SHAP for explainability
import shap

# Utilities
import pickle
import warnings
warnings.filterwarnings('ignore')

# Set visualization style
plt.style.use('seaborn-v0_8-whitegrid')
sns.set_palette("husl")

print("✅ All libraries imported successfully!")

✅ All libraries imported successfully!
```

2. Load Dataset

```
# =====
# STEP 2: LOAD THE DATASET
# =====

df = pd.read_csv('WA_Fn-UseC_-HR-Employee-Attrition.csv')

print("*"*50)
print("DATASET LOADED SUCCESSFULLY")
print("*"*50)
print(f"\nShape: {df.shape[0]} rows x {df.shape[1]} columns")
print(f"\nFirst 5 rows:")
df.head()
```

```
=====
DATASET LOADED SUCCESSFULLY
=====
```

Shape: 1470 rows × 35 columns

First 5 rows:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber	...	RelationshipS...
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sciences	1	1	1	...
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Sciences	1	2	2	...
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	Other	1	4	4	...
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	1	5	5	...
4	27	No	Travel_Rarely	591	Research & Development	2	1	Medical	1	7	7	...

5 rows × 35 columns

3. Data Exploration

```
# =====
# STEP 3: EXPLORE THE DATA
# =====

print("*"*50)
print("DATA EXPLORATION")
print("*"*50)

# Column names
print(f"\nColumns ({len(df.columns)}):")
print(list(df.columns))

# Data types
print("\nData Types:")
print(df.dtypes)

# Missing values
print("\nMissing Values:")
missing = df.isnull().sum()
print(missing[missing > 0] if missing.sum() > 0 else "No missing values found!")

# Duplicates
print(f"\nDuplicate Rows: {df.duplicated().sum()}")

=====
DATA EXPLORATION
=====

Columns (35):
['Age', 'Attrition', 'BusinessTravel', 'DailyRate', 'Department', 'DistanceFromHome', 'Education', 'EducationField', 'EmployeeCount', 'EmployeeNumber', 'EnvironmentSatisfaction', 'Gender', 'HourlyRate', 'JobInvolvement', 'JobLevel', 'JobRole', 'JobSatisfaction', 'MaritalStatus', 'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked', 'Over18', 'Overtime', 'PercentSalaryHike', 'PerformanceRating', 'RelationshipSatisfaction', 'StandardHours', 'StockOptionLevel', 'TotalWorkingYears', 'TrainingTimesLastYear', 'WorkLifeBalance', 'YearsAtCompany', 'YearsInCurrentRole', 'YearsSinceLastPromotion', 'YearsWithCurrManager']

Data Types:
Age          int64
Attrition    object
BusinessTravel  object
DailyRate     int64
Department    object
DistanceFromHome  int64
Education     int64
EducationField  object
EmployeeCount  int64
EmployeeNumber  int64
EnvironmentSatisfaction  int64
Gender        object
HourlyRate     int64
JobInvolvement  int64
JobLevel       int64
JobRole        object
JobSatisfaction  int64
MaritalStatus   object
MonthlyIncome   int64
MonthlyRate     int64
NumCompaniesWorked  int64
Over18         object
Overtime        object
PercentSalaryHike  int64
PerformanceRating  int64
RelationshipSatisfaction  int64
StandardHours   int64
StockOptionLevel  int64
TotalWorkingYears  int64
TrainingTimesLastYear  int64
WorkLifeBalance  int64
YearsAtCompany   int64
YearsInCurrentRole  int64
YearsSinceLastPromotion  int64
YearsWithCurrManager  int64
dtype: object

Missing Values:
No missing values found!

Duplicate Rows: 0
```

```
# Check target variable distribution
print("\n" + "*50)
print("TARGET VARIABLE (ATTRITION) DISTRIBUTION")
print("*50)

attrition_counts = df['Attrition'].value_counts()
print(f"\n{attrition_counts}")
print(f"\nAttrition Rate: {attrition_counts['Yes'] / len(df) * 100:.1f}%")

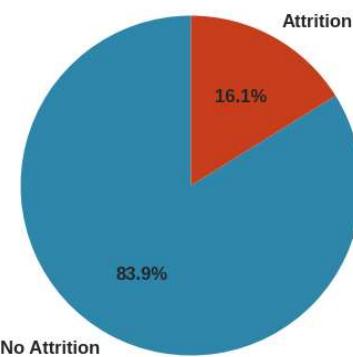
# Plot
plt.figure(figsize=(8, 5))
colors = ['#2E86A8', '#C73E1D']
plt.pie(attrition_counts, labels=['No Attrition', 'Attrition'],
        autopct='%.1f%%', colors=colors, startangle=90,
        textprops={'fontsize': 12, 'fontweight': 'bold'})
plt.title('Overall Attrition Distribution', fontsize=14, fontweight='bold')
plt.show()
```

```
=====
TARGET VARIABLE (ATTRITION) DISTRIBUTION
=====
```

```
Attrition
No      1233
Yes     237
Name: count, dtype: int64
```

```
Attrition Rate: 16.1%
```

Overall Attrition Distribution



4. Data Preprocessing

```
# =====
# STEP 4: DATA PREPROCESSING
# =====

print("*50)
print("DATA PREPROCESSING")
print("*50)

# Create a copy for preprocessing
df_processed = df.copy()

# 4.1: Encode target variable (Attrition)
df_processed['Attrition'] = df_processed['Attrition'].map({'Yes': 1, 'No': 0})
print("\n1. Target variable encoded: Yes=1, No=0")

# 4.2: Remove unnecessary columns
columns_to_drop = ['EmployeeCount', 'EmployeeNumber', 'Over18', 'StandardHours']
df_processed = df_processed.drop(columns=columns_to_drop, errors='ignore')
print(f"2. Dropped columns: {columns_to_drop}")

# 4.3: Encode categorical variables
categorical_cols = df_processed.select_dtypes(include=['object']).columns.tolist()
if 'Attrition' in categorical_cols:
    categorical_cols.remove('Attrition')

label_encoders = {}
for col in categorical_cols:
    le = LabelEncoder()
    df_processed[col] = le.fit_transform(df_processed[col].astype(str))
    label_encoders[col] = le

print(f"3. Encoded {len(categorical_cols)} categorical columns")
print(f"  Columns: {categorical_cols}")

print(f"\nProcessed dataset shape: {df_processed.shape}")
df_processed.head()
```

```
=====
DATA PREPROCESSING
=====

1. Target variable encoded: Yes=1, No=0
2. Dropped columns: ['EmployeeCount', 'EmployeeNumber', 'Over18', 'StandardHours']
3. Encoded 7 categorical columns
    Columns: ['BusinessTravel', 'Department', 'EducationField', 'Gender', 'JobRole', 'MaritalStatus', 'Overtime']

Processed dataset shape: (1470, 31)

   Age Attrition BusinessTravel DailyRate Department DistanceFromHome Education EducationField EnvironmentSatisfaction Gender ... PerformanceRate
0   41        1            2       1102          2              1         2             1                  2           0           ...
1   49        0            1       279           1              8         1             1                  3           1           ...
2   37        1            2      1373           1              2         2             4                  4           1           ...
3   33        0            1      1392           1              3         4             1                  4           0           ...
4   27        0            2       591           1              2         1             3                  1           1           ...

5 rows × 31 columns
```

5. Feature Engineering

```
# =====
# STEP 5: FEATURE ENGINEERING
# =====

print("*"*50)
print("FEATURE ENGINEERING")
print("*"*50)

# 5.1: Create derived features

# Average years per company
df_processed['AverageYearsPerCompany'] = df_processed['TotalWorkingYears'] / (
    df_processed['NumCompaniesWorked'] + 1
)

# Years since last promotion ratio
df_processed['YearsSinceLastPromotionRatio'] = df_processed['YearsSinceLastPromotion'] / (
    df_processed['YearsAtCompany'] + 1
)

# Income per year of experience
df_processed['IncomePerYearExperience'] = df_processed['MonthlyIncome'] / (
    df_processed['TotalWorkingYears'] + 1
)

# Age groups
df_processed['AgeGroup'] = pd.cut(
    df_processed['Age'],
    bins=[0, 30, 40, 50, 100],
    labels=[0, 1, 2, 3]
).astype(int)

# Overall satisfaction score
df_processed['OverallSatisfaction'] = (
    df_processed['JobSatisfaction'] +
    df_processed['EnvironmentSatisfaction'] +
    df_processed['WorkLifeBalance']
) / 3

print("Created 5 new features:")
print("  1. AverageYearsPerCompany")
print("  2. YearsSinceLastPromotionRatio")
print("  3. IncomePerYearExperience")
print("  4. AgeGroup")
print("  5. OverallSatisfaction")

print(f"\nFinal dataset shape: {df_processed.shape}")

=====

FEATURE ENGINEERING
=====

Created 5 new features:
  1. AverageYearsPerCompany
  2. YearsSinceLastPromotionRatio
  3. IncomePerYearExperience
  4. AgeGroup
  5. OverallSatisfaction

Final dataset shape: (1470, 36)
```

6. Train-Test Split

```
# =====
# STEP 6: SPLIT DATA INTO TRAIN AND TEST
# =====

print("*"*50)
print("TRAIN-TEST SPLIT")
print("*"*50)
```

```
# Separate features and target
X = df_processed.drop('Attrition', axis=1)
y = df_processed['Attrition']

# Split data
X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.2,      # 80% train, 20% test
    random_state=42,    # For reproducibility
    stratify=y          # Maintain class distribution
)

print(f"\nTraining set: {X_train.shape[0]} samples")
print(f"Test set: {X_test.shape[0]} samples")

# Feature scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Convert back to DataFrame
X_train_scaled = pd.DataFrame(X_train_scaled, columns=X.columns, index=X_train.index)
X_test_scaled = pd.DataFrame(X_test_scaled, columns=X.columns, index=X_test.index)

print("\nFeatures scaled using StandardScaler")
print("\n

```

```
# 7.2: Decision Tree
print("\n2. Training Decision Tree...")
dt = DecisionTreeClassifier(
    max_depth=10,
    min_samples_split=5,
    random_state=42,
    class_weight='balanced'
)
dt.fit(X_train_scaled, y_train)
models['Decision Tree'] = dt

cv_scores = cross_val_score(dt, X_train_scaled, y_train, cv=5, scoring='roc_auc')
results['Decision Tree'] = {'cv_score': cv_scores.mean()}
print(f" CV ROC-AUC: {cv_scores.mean():.4f}")

=====

2. Training Decision Tree...
CV ROC-AUC: 0.6118
```

```
# 7.3: Random Forest
print("\n3. Training Random Forest...")
rf = RandomForestClassifier(
    n_estimators=100,
    max_depth=15,
    random_state=42,
    class_weight='balanced'
```

```

)
rf.fit(X_train_scaled, y_train)
models['Random Forest'] = rf

cv_scores = cross_val_score(rf, X_train_scaled, y_train, cv=5, scoring='roc_auc')
results['Random Forest'] = {'cv_score': cv_scores.mean()}
print("CV ROC AUC: ", cv_scores.mean(), "(+/-", 4*cv_scores.std(), ")")

```

3. Training Random Forest...
CV ROC-AUC: 0.7917

```
# 7.4: Gradient Boosting
print("\n4. Training Gradient Boosting...")
gb = GradientBoostingClassifier(
    n_estimators=100,
    learning_rate=0.1,
    max_depth=5,
    random_state=42
)
gb.fit(X_train_scaled, y_train)
models['Gradient Boosting'] = gb
```

```
cv_scores = cross_val_score(gb, X_train_scaled, y_train, cv=5, scoring='roc_auc')
results['Gradient Boosting'] = {'cv_score': cv_scores.mean()}
print(f"  CV ROC-AUC: {cv_scores.mean():.4f}")
```

CV ROC-AUC: 0.7730

```

# 7.5: XGBoost (Best Model)
print("\n5. Training XGBoost...")
xgb_model = xgb.XGBClassifier(
    learning_rate=0.1,
    max_depth=6,
    n_estimators=200,
    subsample=0.8,
    colsample_bytree=0.8,
    random_state=42,
    eval_metric='logloss'
)
xgb_model.fit(X_train_scaled, y_train)
models['XGBoost'] = xgb_model

cv_scores = cross_val_score(xgb_model, X_train_scaled, y_train, cv=5, scoring='roc_auc')
results['XGBoost'] = {'cv_score': cv_scores.mean()}
print(f" CV ROC-AUC: {cv_scores.mean():.4f}")

```

5. Training XGBoost...
CV ROC-AUC: 0.8008

```
# Print training summary
print("\n" + "*50)
print("TRAINING SUMMARY")
print("*50)

for name, result in results.items():
    print(f"{name}: {result['cv_score']:.4f}")

# Find best model
best_model_name = max(results, key=lambda x: results[x]['cv_score'])
best_model = models[best_model_name]

print(f"\n✓ Best Model: {best_model_name}")
print(f"CV ROC-AUC: {results[best_model_name]['cv_score']:.4f}")
```

```
=====
TRAINING SUMMARY
=====
Logistic Regression: 0.8042
Decision Tree: 0.6118
Random Forest: 0.7917
Gradient Boosting: 0.7730
XGBoost: 0.8008
```

 Best Model: Logistic Regression
CV ROC-AUC: 0.8042

▼ 8. Model Evaluation

```
# =====
# STEP 8: EVALUATE ALL MODELS
# =====

print("=*50)
print("MODEL EVALUATION")
print("=*50)

evaluation_results = []

for name, model in models.items():
    # Make predictions
    y_pred = model.predict(X_test_scaled)
    y_pred_proba = model.predict_proba(y_test)
```

```
# Calculate metrics
metrics = {
    'Model': name,
    'Accuracy': accuracy_score(y_test, y_pred),
    'Precision': precision_score(y_test, y_pred),
    'Recall': recall_score(y_test, y_pred),
    'F1-Score': f1_score(y_test, y_pred),
    'ROC-AUC': roc_auc_score(y_test, y_pred_proba)
}

evaluation_results.append(metrics)

# Create comparison DataFrame
results_df = pd.DataFrame(evaluation_results)

print("\nModel Performance Comparison:")
print(results_df.to_string(index=False))

=====
MODEL EVALUATION
=====
```

Model Performance Comparison:

Model	Accuracy	Precision	Recall	F1-Score	ROC-AUC
Logistic Regression	0.748299	0.363636	0.765957	0.493151	0.802998
Decision Tree	0.785714	0.326087	0.319149	0.322581	0.622104
Random Forest	0.833333	0.375000	0.063830	0.109091	0.751400
Gradient Boosting	0.836735	0.476190	0.212766	0.294118	0.750188
XGBoost	0.857143	0.666667	0.212766	0.322581	0.741752

```
# Detailed evaluation of best model
print("\n" + "="*50)
print(f"DETAILED EVALUATION: {best_model_name}")
print("="*50)

y_pred_best = best_model.predict(X_test_scaled)
y_pred_proba_best = best_model.predict_proba(X_test_scaled)[:, 1]

print("\nClassification Report:")
print(classification_report(y_test, y_pred_best,
                           target_names=['No Attrition', 'Attrition']))
```

```
=====
DETAILED EVALUATION: Logistic Regression
=====

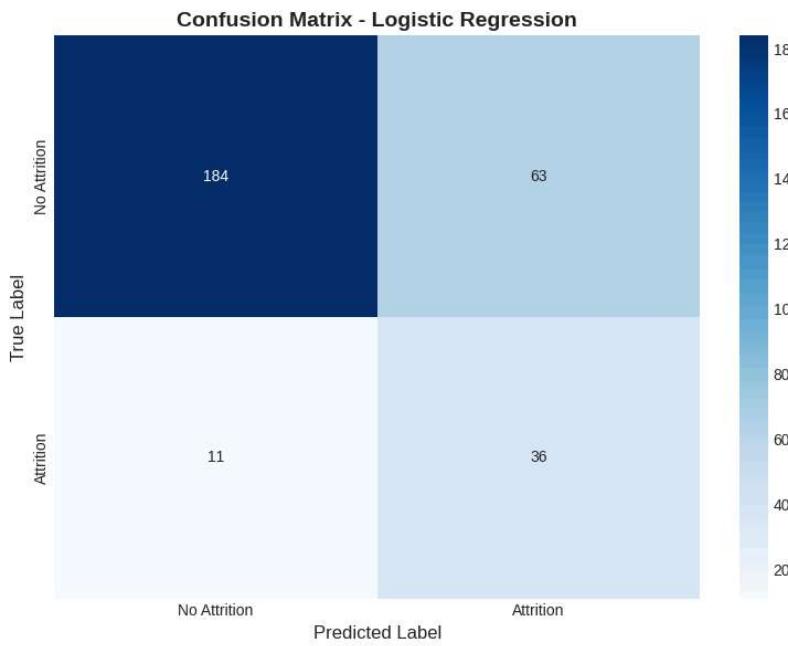
Classification Report:
precision    recall   f1-score   support
No Attrition    0.94     0.74     0.83      247
Attrition       0.36     0.77     0.49       47

accuracy          0.75
macro avg        0.65     0.76     0.66      294
weighted avg     0.85     0.75     0.78      294
```

```
# Confusion Matrix
cm = confusion_matrix(y_test, y_pred_best)

plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['No Attrition', 'Attrition'],
            yticklabels=['No Attrition', 'Attrition'])
plt.title(f'Confusion Matrix - {best_model_name}', fontsize=14, fontweight='bold')
plt.ylabel('True Label', fontsize=12)
plt.xlabel('Predicted Label', fontsize=12)
plt.tight_layout()
plt.show()

print("\nConfusion Matrix:")
print(f" True Negatives: {cm[0,0]}")
print(f" False Positives: {cm[0,1]}")
print(f" False Negatives: {cm[1,0]}")
print(f" True Positives: {cm[1,1]}")
```



```
Confusion Matrix:
  True Negatives: 184
  False Positives: 63
  False Negatives: 11
  True Positives: 36
```

9. SHAP Explainability

```
# =====
# STEP 9: SHAP EXPLAINABILITY
# =====

print("*"*50)
print("SHAP EXPLAINABILITY ANALYSIS")
print("*"*50)

import shap

# Initialize SHAP explainer for linear models
explainer = shap.LinearExplainer(best_model, X_train_scaled)

# Calculate SHAP values
shap_values = explainer.shap_values(X_test_scaled)

# For binary classification, use values for class 1
if isinstance(shap_values, list):
    shap_values = shap_values[1]

print("✅ SHAP values calculated!")

=====
```

```
SHAP EXPLAINABILITY ANALYSIS
=====
✅ SHAP values calculated!
```

```
# Feature Importance (Mean Absolute SHAP values)
importance = np.abs(shap_values).mean(axis=0)

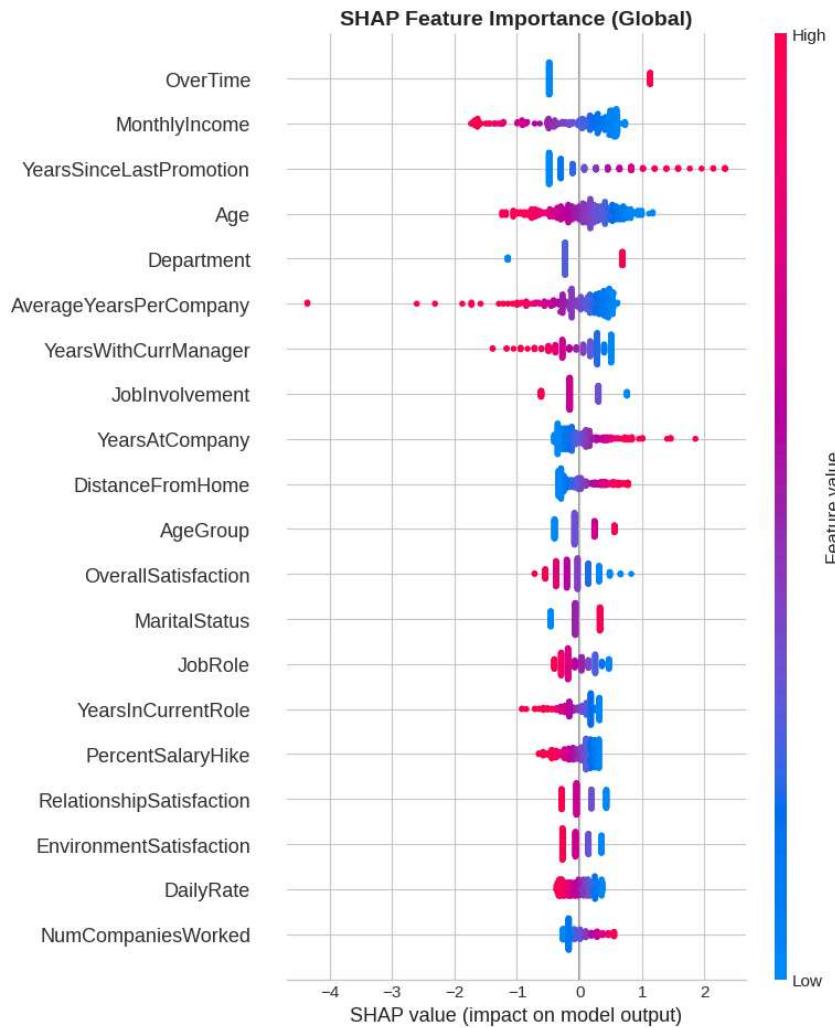
feature_importance = pd.DataFrame({
    'Feature': X.columns,
    'Importance': importance
}).sort_values('Importance', ascending=False)

print("\nTop 10 Most Important Features:")
print(feature_importance.head(10).to_string(index=False))
```

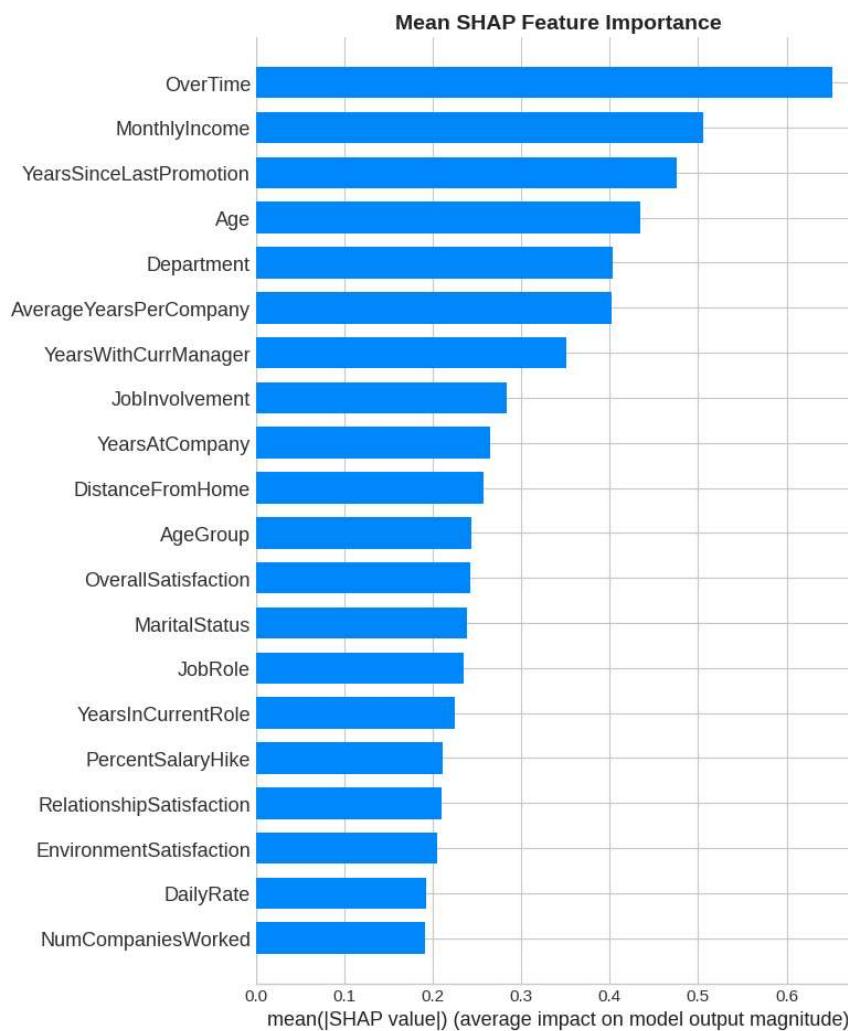
```
Top 10 Most Important Features:
      Feature  Importance
      Overtime   0.651626
      MonthlyIncome  0.595159
      YearsSinceLastPromotion  0.475348
      Age  0.434405
      Department  0.403781
      AverageYearsPerCompany  0.401640
      YearsWithCurrManager  0.350982
      JobInvolvement  0.283521
      YearsAtCompany  0.265264
      DistanceFromHome  0.257136
```

```
# SHAP Summary Plot (Global Feature Importance)
plt.figure(figsize=(10, 8))
shap.summary_plot(shap_values, X_test_scaled, show=False)
plt.title('SHAP Feature Importance (Global)', fontsize=14, fontweight='bold')
```

```
plt.tight_layout()
plt.show()
```



```
# SHAP Bar Plot (Mean Absolute Importance)
plt.figure(figsize=(10, 8))
shap.summary_plot(shap_values, X_test_scaled, plot_type="bar", show=False)
plt.title('Mean SHAP Feature Importance', fontsize=14, fontweight='bold')
plt.tight_layout()
plt.show()
```

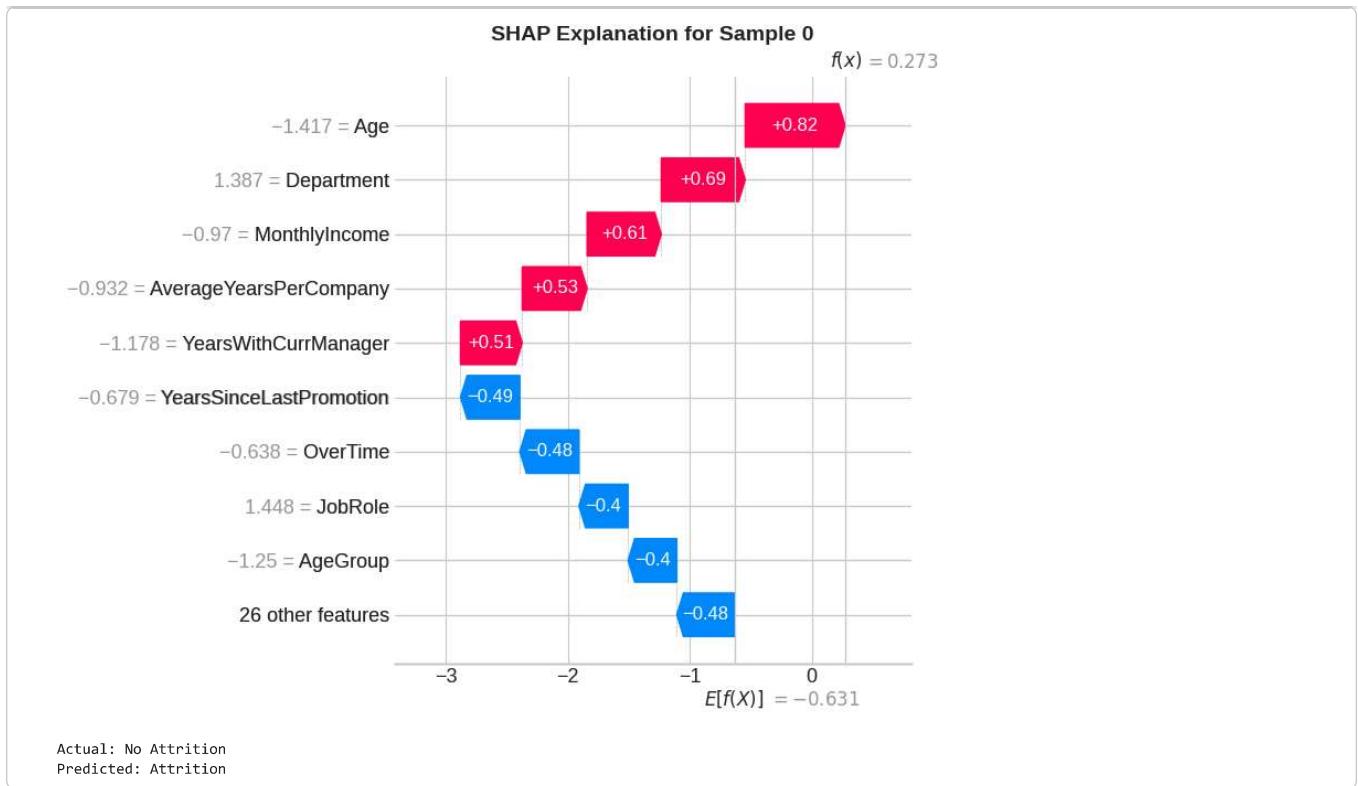


```
# Explain a single prediction (Waterfall Plot)
sample_idx = 0 # Change this to explore different samples

explanation = shap.Explanation(
    values=shap_values[sample_idx],
    base_values=explainer.expected_value,
    data=X_test_scaled.iloc[sample_idx],
    feature_names=X.columns
)

plt.figure(figsize=(12, 8))
shap.waterfall_plot(explanation, show=False)
plt.title(f'SHAP Explanation for Sample {sample_idx}', fontsize=14, fontweight='bold')
plt.tight_layout()
plt.show()

actual = y_test.iloc[sample_idx]
predicted = y_pred_best[sample_idx]
print(f"\nActual: {'Attrition' if actual == 1 else 'No Attrition'}")
print(f"Predicted: {'Attrition' if predicted == 1 else 'No Attrition'}")
```



10. Visualization

```
# =====
# STEP 10: CREATE VISUALIZATIONS
# =====

print("*"*50)
print("CREATING VISUALIZATIONS")
print("*"*50)

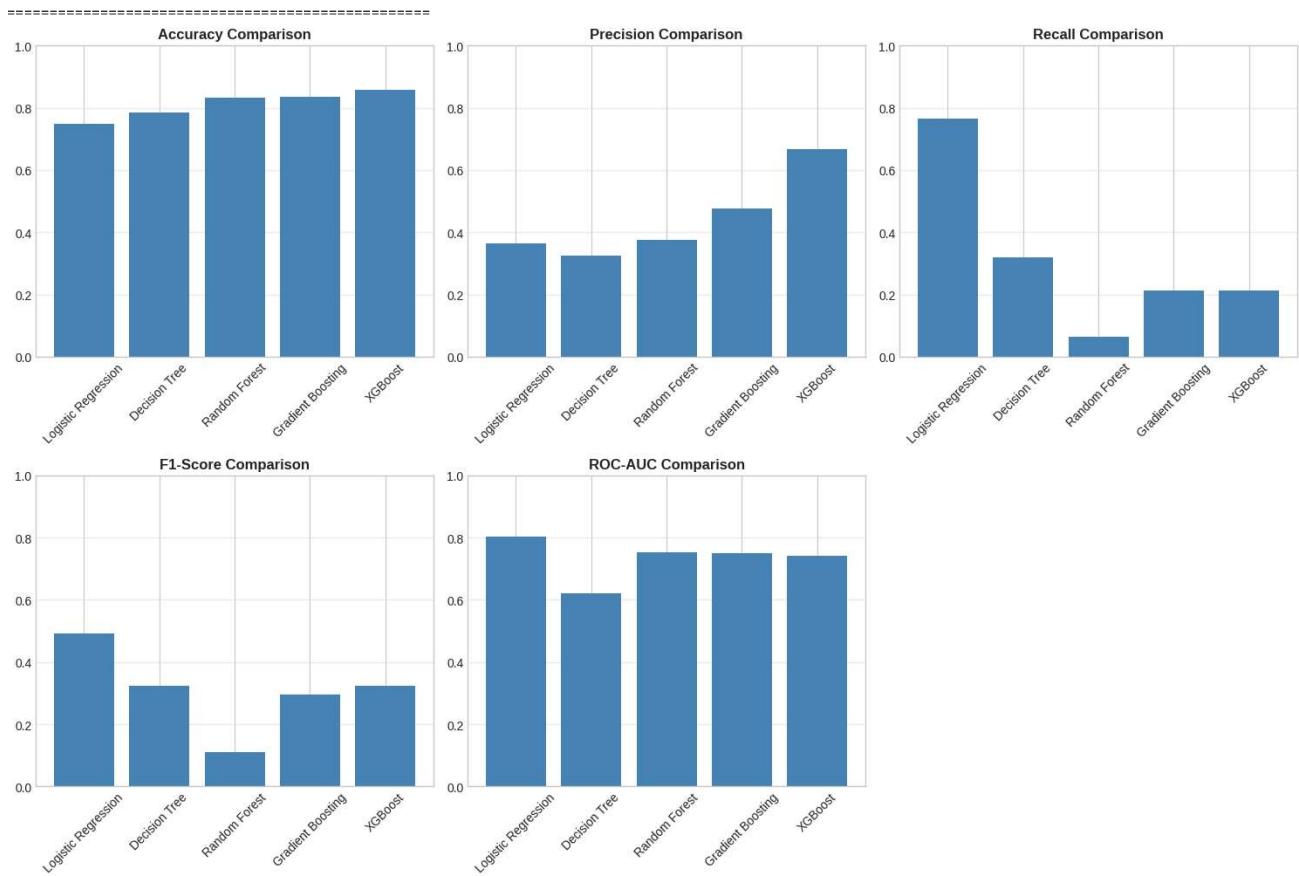
# 10.1: Model Performance Comparison
fig, axes = plt.subplots(2, 3, figsize=(15, 10))
axes = axes.flatten()

metrics = ['Accuracy', 'Precision', 'Recall', 'F1-Score', 'ROC-AUC']

for i, metric in enumerate(metrics):
    axes[i].bar(results_df['Model'], results_df[metric], color='steelblue')
    axes[i].set_title(f'{metric} Comparison', fontsize=12, fontweight='bold')
    axes[i].set_ylim([0, 1])
    axes[i].tick_params(axis='x', rotation=45)
    axes[i].grid(axis='y', alpha=0.3)

axes[5].axis('off')
plt.tight_layout()
plt.show()
```

```
=====
CREATING VISUALIZATIONS
=====
```

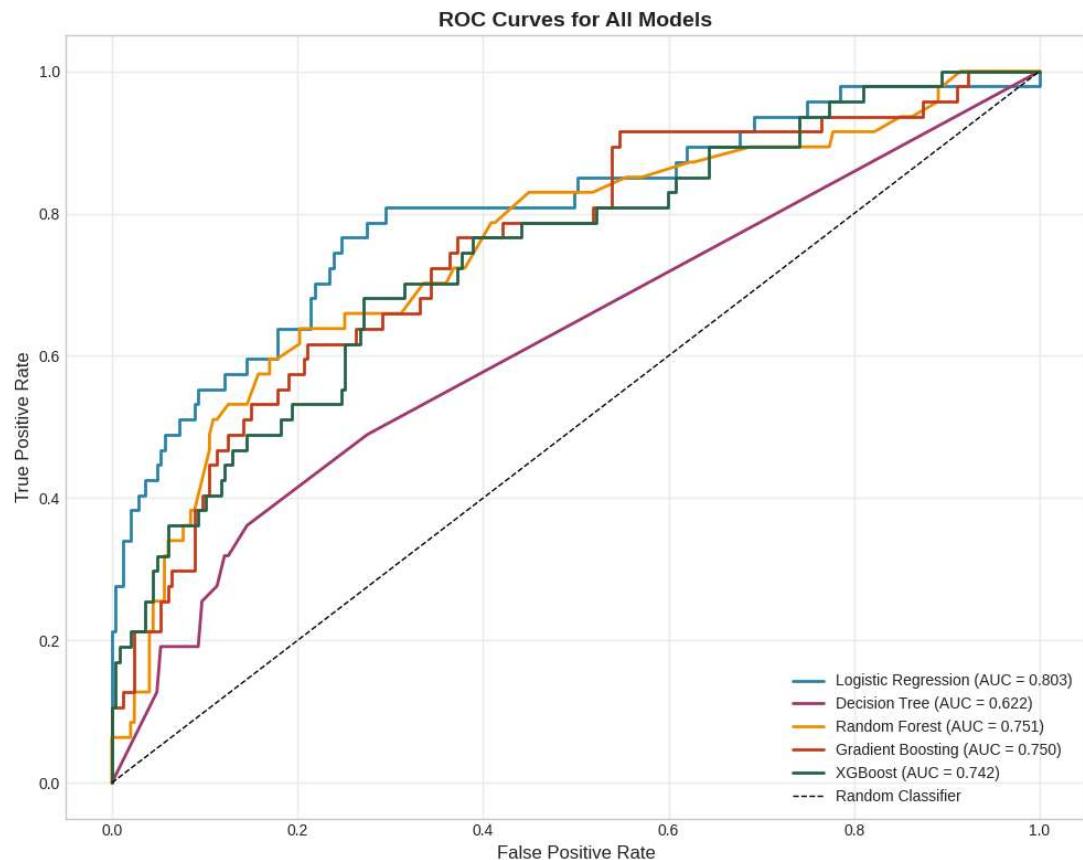


```
# 10.2: ROC Curves for All Models
plt.figure(figsize=(10, 8))

colors = ['#2E86AB', '#A23B72', '#F18F01', '#C73E1D', '#276749']

for i, (name, model) in enumerate(models.items()):
    y_pred_proba = model.predict_proba(X_test_scaled)[:, 1]
    fpr, tpr, _ = roc_curve(y_test, y_pred_proba)
    auc = roc_auc_score(y_test, y_pred_proba)
    plt.plot(fpr, tpr, linewidth=2, label=f'{name} (AUC = {auc:.3f})', color=colors[i])

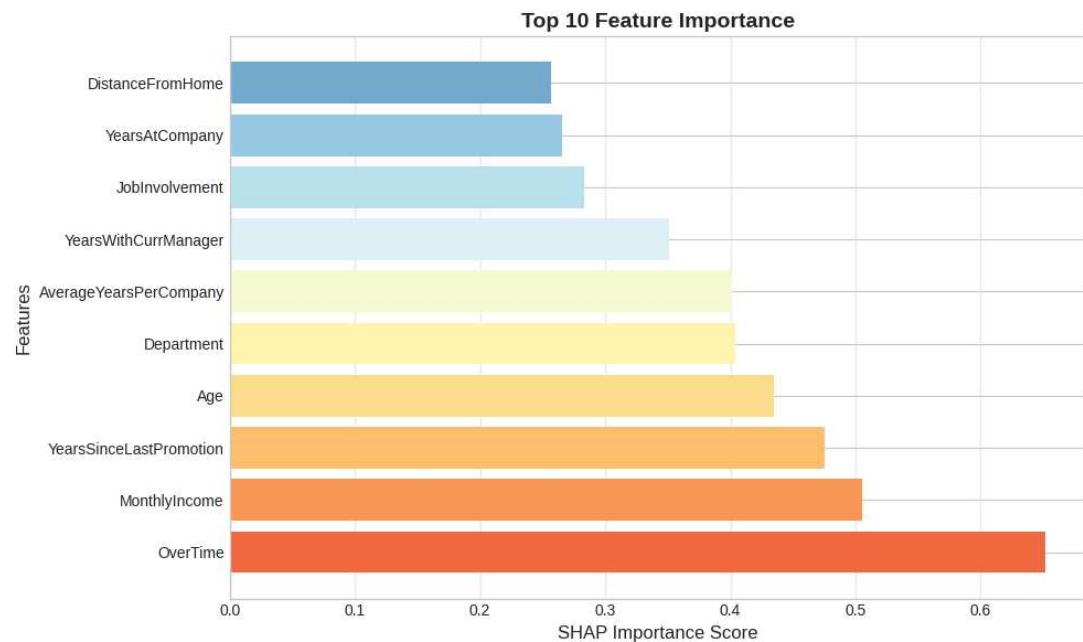
plt.plot([0, 1], [0, 1], 'k--', linewidth=1, label='Random Classifier')
plt.xlabel('False Positive Rate', fontsize=12)
plt.ylabel('True Positive Rate', fontsize=12)
plt.title('ROC Curves for All Models', fontsize=14, fontweight='bold')
plt.legend(loc='lower right')
plt.grid(alpha=0.3)
plt.tight_layout()
plt.show()
```



```
# 10.3: Feature Importance Bar Chart
plt.figure(figsize=(10, 6))

top_features = feature_importance.head(10)
colors = plt.cm.RdYlBu(np.linspace(0.2, 0.8, len(top_features)))

plt.barh(top_features['Feature'], top_features['Importance'], color=colors)
plt.xlabel('SHAP Importance Score', fontsize=12)
plt.ylabel('Features', fontsize=12)
plt.title('Top 10 Feature Importance', fontsize=14, fontweight='bold')
plt.grid(axis='x', alpha=0.3)
plt.tight_layout()
plt.show()
```



11. Save Model

```
# =====
# STEP 11: SAVE MODEL AND PREPROCESSOR
# =====

import os
```

```
# Create models directory
os.makedirs('models', exist_ok=True)

# Save best model
with open('models/xgboost_model.pkl', 'wb') as f:
    pickle.dump(best_model, f)

# Save scaler
with open('models/scaler.pkl', 'wb') as f:
    pickle.dump(scaler, f)

# Save label encoders
with open('models/label_encoders.pkl', 'wb') as f:
    pickle.dump(label_encoders, f)

# Save feature columns
with open('models/feature_columns.pkl', 'wb') as f:
    pickle.dump(list(X.columns), f)

print("=*50")
print("MODELS SAVED SUCCESSFULLY")
print("=*50")
print("\nSaved files:")
print(" - models/xgboost_model.pkl")
print(" - models/scaler.pkl")
print(" - models/label_encoders.pkl")
print(" - models/feature_columns.pkl")

=====
===== MODELS SAVED SUCCESSFULLY =====
=====
```

Saved files: