

Employee Attrition Prediction

Student: Amruta Kumbar

Roll Number: AA.SC.P2MCA24074031

Course: MCA Minor Project (21CSA697A)

Date: February 2026

This notebook contains the complete code for Employee Attrition Prediction using Machine Learning and Explainable AI.

Table of Contents

1. [Import Libraries](#)
2. [Load Dataset](#)
3. [Data Exploration](#)
4. [Data Preprocessing](#)
5. [Feature Engineering](#)
6. [Train-Test Split](#)
7. [Model Training](#)
8. [Model Evaluation](#)
9. [SHAP Explainability](#)
10. [Visualization](#)
11. [Save Model](#)
12. [Make Predictions](#)

1. Import Libraries

```
# =====
# STEP 1: IMPORT ALL REQUIRED LIBRARIES
# =====

# Data manipulation
import pandas as pd
import numpy as np

# Visualization
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Machine Learning
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    roc_auc_score, confusion_matrix, classification_report, roc_curve
)

# XGBoost
import xgboost as xgb

# SHAP for explainability
import shap

# Utilities
import pickle
import warnings
warnings.filterwarnings('ignore')

# Set visualization style
plt.style.use('seaborn-v0_8-whitegrid')
sns.set_palette("husl")

print("✅ All libraries imported successfully!")
```

✅ All libraries imported successfully!

2. Load Dataset

```
# =====
# STEP 2: LOAD THE DATASET
# =====

df = pd.read_csv('WA_Fn-UseC_-HR-Employee-Attrition.csv')

print("*50)
print("DATASET LOADED SUCCESSFULLY")
print("*50)
print(f"\nShape: {df.shape[0]} rows x {df.shape[1]} columns")
print(f"\nFirst 5 rows:")
df.head()
```

```
=====
DATASET LOADED SUCCESSFULLY
=====
```

Shape: 1470 rows × 35 columns

First 5 rows:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Edu
0	41	Yes	Travel_Rarely	1102	Sales		1
1	49	No	Travel_Frequently	279	Research & Development		8
2	37	Yes	Travel_Rarely	1373	Research & Development		2
3	33	No	Travel_Frequently	1392	Research & Development		3
4	27	No	Travel_Rarely	591	Research & Development		2

5 rows × 35 columns

3. Data Exploration

```
# =====
# STEP 3: EXPLORE THE DATA
# =====

print("*50)
print("DATA EXPLORATION")
print("*50)

# Column names
print(f"\nColumns ({len(df.columns)}):")
print(list(df.columns))

# Data types
print("\nData Types:")
print(df.dtypes)

# Missing values
print("\nMissing Values:")
missing = df.isnull().sum()
print(missing[missing > 0] if missing.sum() > 0 else "No missing values found!")
```

```
# Duplicates
print(f"\nDuplicate Rows: {df.duplicated().sum()}")


=====
DATA EXPLORATION
=====

Columns (35):
['Age', 'Attrition', 'BusinessTravel', 'DailyRate', 'Department', 'DistanceFromHome', 'Education', 'EducationField', 'EmployeeCount', 'EmployeeNumber', 'EnvironmentSatisfaction', 'Gender', 'HourlyRate', 'JobInvolvement', 'JobLevel', 'JobRole', 'JobSatisfaction', 'MaritalStatus', 'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked', 'Over18', 'Overtime', 'PercentSalaryHike', 'PerformanceRating', 'RelationshipSatisfaction', 'StandardHours', 'StockOptionLevel', 'TotalWorkingYears', 'TrainingTimesLastYear', 'WorkLifeBalance', 'YearsAtCompany', 'YearsInCurrentRole', 'YearsSinceLastPromotion', 'YearsWithCurrManager']

Data Types:
Age                      int64
Attrition                object
BusinessTravel            object
DailyRate                 int64
Department               object
DistanceFromHome          int64
Education                int64
EducationField            object
EmployeeCount             int64
EmployeeNumber            int64
EnvironmentSatisfaction int64
Gender                   object
HourlyRate                int64
JobInvolvement            int64
JobLevel                  int64
JobRole                   object
JobSatisfaction           int64
MaritalStatus              object
MonthlyIncome              int64
MonthlyRate                int64
NumCompaniesWorked         int64
Over18                    object
Overtime                  object
PercentSalaryHike          int64
PerformanceRating          int64
RelationshipSatisfaction int64
StandardHours              int64
StockOptionLevel            int64
TotalWorkingYears           int64
TrainingTimesLastYear       int64
WorkLifeBalance             int64
YearsAtCompany              int64
YearsInCurrentRole          int64
YearsSinceLastPromotion      int64
YearsWithCurrManager         int64
dtype: object

Missing Values:
No missing values found!

Duplicate Rows: 0
```

```
# Check target variable distribution
print("\n" + "*50)
print("TARGET VARIABLE (ATTRITION) DISTRIBUTION")
print("*50)

attrition_counts = df['Attrition'].value_counts()
print(f"\n{attrition_counts}")
print(f"\nAttrition Rate: {attrition_counts['Yes'] / len(df) * 100:.1f}%")

# Plot
plt.figure(figsize=(8, 5))
colors = ['#2E86AB', '#C73E1D']
plt.pie(attrition_counts, labels=['No Attrition', 'Attrition'],
        autopct='%.1f%%', colors=colors, startangle=90,
        textprops={'fontsize': 12, 'fontweight': 'bold'})
plt.title('Overall Attrition Distribution', fontsize=14, fontweight='bold')
plt.show()
```

```
=====
```

```
TARGET VARIABLE (ATTRITION) DISTRIBUTION
```

```
=====
```

```
Attrition
```

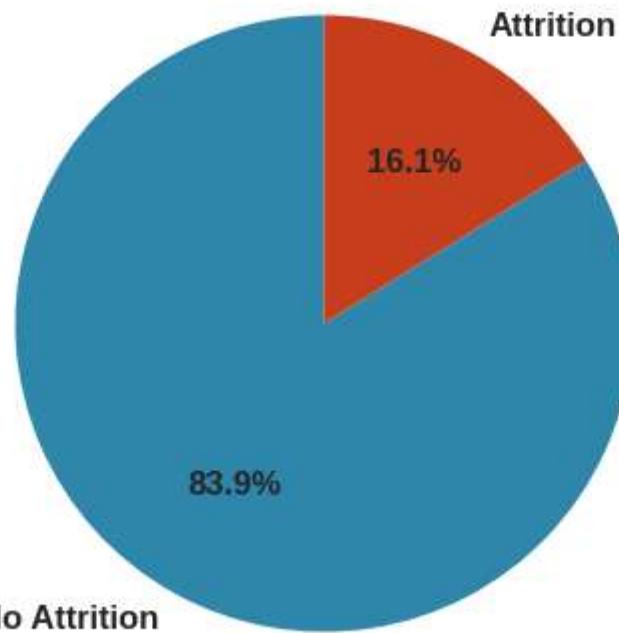
```
No      1233
```

```
Yes     237
```

```
Name: count, dtype: int64
```

```
Attrition Rate: 16.1%
```

Overall Attrition Distribution



▼ 4. Data Preprocessing

```
# =====
# STEP 4: DATA PREPROCESSING
# =====

print("*50)
print("DATA PREPROCESSING")
print("*50)

# Create a copy for preprocessing
df_processed = df.copy()

# 4.1: Encode target variable (Attrition)
df_processed['Attrition'] = df_processed['Attrition'].map({'Yes': 1, 'No': 0})
print("\n1. Target variable encoded: Yes=1, No=0")

# 4.2: Remove unnecessary columns
columns_to_drop = ['EmployeeCount', 'EmployeeNumber', 'Over18', 'StandardHours']
df_processed = df_processed.drop(columns=columns_to_drop, errors='ignore')
print(f"2. Dropped columns: {columns_to_drop}")

# 4.3: Encode categorical variables
categorical_cols = df_processed.select_dtypes(include=['object']).columns.tolist()
if 'Attrition' in categorical_cols:
    categorical_cols.remove('Attrition')

label_encoders = {}
for col in categorical_cols:
    le = LabelEncoder()
    df_processed[col] = le.fit_transform(df_processed[col].astype(str))
    label_encoders[col] = le

print(f"3. Encoded {len(categorical_cols)} categorical columns")
print(f"    Columns: {categorical_cols}")

print(f"\nProcessed dataset shape: {df_processed.shape}")
df_processed.head()
```

```
=====
DATA PREPROCESSING
=====
```

1. Target variable encoded: Yes=1, No=0
2. Dropped columns: ['EmployeeCount', 'EmployeeNumber', 'Over18', 'StandardHours']
3. Encoded 7 categorical columns
Columns: ['BusinessTravel', 'Department', 'EducationField', 'Gender', 'JobRole', 'MaritalStatus', 'Relationship']

Processed dataset shape: (1470, 31)

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	EducationField
0	41	1	2	1102	2		1
1	49	0	1	279	1		8
2	37	1	2	1373	1		2
3	33	0	1	1392	1		3
4	27	0	2	591	1		2

5 rows × 31 columns

▼ 5. Feature Engineering

```
# =====
# STEP 5: FEATURE ENGINEERING
# =====

print("*"*50)
print("FEATURE ENGINEERING")
print("*"*50)

# 5.1: Create derived features

# Average years per company
df_processed['AverageYearsPerCompany'] = df_processed['TotalWorkingYears'] / (
    df_processed['NumCompaniesWorked'] + 1
)

# Years since last promotion ratio
df_processed['YearsSinceLastPromotionRatio'] = df_processed['YearsSinceLastPromotion'] / (
    df_processed['YearsAtCompany'] + 1
)

# Income per year of experience
df_processed['IncomePerYearExperience'] = df_processed['MonthlyIncome'] / (
    df_processed['TotalWorkingYears'] + 1
)
```

```

)
# Age groups
df_processed[ 'AgeGroup' ] = pd.cut(
    df_processed[ 'Age' ],
    bins=[0, 30, 40, 50, 100],
    labels=[0, 1, 2, 3]
).astype(int)

# Overall satisfaction score
df_processed[ 'OverallSatisfaction' ] = (
    df_processed[ 'JobSatisfaction' ] +
    df_processed[ 'EnvironmentSatisfaction' ] +
    df_processed[ 'WorkLifeBalance' ]
) / 3

print("Created 5 new features:")
print(" 1. AverageYearsPerCompany")
print(" 2. YearsSinceLastPromotionRatio")
print(" 3. IncomePerYearExperience")
print(" 4. AgeGroup")
print(" 5. OverallSatisfaction")

print(f"\nFinal dataset shape: {df_processed.shape}")

```

```

=====
FEATURE ENGINEERING
=====
Created 5 new features:
 1. AverageYearsPerCompany
 2. YearsSinceLastPromotionRatio
 3. IncomePerYearExperience
 4. AgeGroup
 5. OverallSatisfaction

```

Final dataset shape: (1470, 36)

6. Train-Test Split

```

# =====
# STEP 6: SPLIT DATA INTO TRAIN AND TEST
# =====

print("*"*50)
print("TRAIN-TEST SPLIT")
print("*"*50)

# Separate features and target
X = df_processed.drop('Attrition', axis=1)

```

```
y = df_processed['Attrition']

# Split data
X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.2,      # 80% train, 20% test
    random_state=42,    # For reproducibility
    stratify=y          # Maintain class distribution
)

print(f"\nTraining set: {X_train.shape[0]} samples")
print(f"Test set: {X_test.shape[0]} samples")

# Feature scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Convert back to DataFrame
X_train_scaled = pd.DataFrame(X_train_scaled, columns=X.columns, index=X_train.index)
X_test_scaled = pd.DataFrame(X_test_scaled, columns=X.columns, index=X_test.index)

print("\nFeatures scaled using StandardScaler")
print("\n✓ Data ready for model training!")
```

```
=====
TRAIN-TEST SPLIT
=====
```

```
Training set: 1176 samples
Test set: 294 samples
```

```
Features scaled using StandardScaler
```

```
✓ Data ready for model training!
```

7. Model Training

```
# =====
# STEP 7: TRAIN MULTIPLE MODELS
# =====

print("*50)
print("MODEL TRAINING")
print("*50)

# Dictionary to store models and results
models = {}
results = {}
```

```
# 7.1: Logistic Regression
print("\n1. Training Logistic Regression...")
lr = LogisticRegression(max_iter=1000, random_state=42, class_weight='balanced')
lr.fit(X_train_scaled, y_train)
models['Logistic Regression'] = lr

# Cross-validation score
cv_scores = cross_val_score(lr, X_train_scaled, y_train, cv=5, scoring='roc_auc')
results['Logistic Regression'] = {'cv_score': cv_scores.mean()}
print(f"    CV ROC-AUC: {cv_scores.mean():.4f}")

=====
```

MODEL TRAINING

```
=====
```

1. Training Logistic Regression...
 CV ROC-AUC: 0.8042

```
# 7.2: Decision Tree
print("\n2. Training Decision Tree...")
dt = DecisionTreeClassifier(
    max_depth=10,
    min_samples_split=5,
    random_state=42,
    class_weight='balanced'
)
dt.fit(X_train_scaled, y_train)
models['Decision Tree'] = dt

cv_scores = cross_val_score(dt, X_train_scaled, y_train, cv=5, scoring='roc_auc')
results['Decision Tree'] = {'cv_score': cv_scores.mean()}
print(f"    CV ROC-AUC: {cv_scores.mean():.4f}")

=====
```

2. Training Decision Tree...
 CV ROC-AUC: 0.6118

```
# 7.3: Random Forest
print("\n3. Training Random Forest...")
rf = RandomForestClassifier(
    n_estimators=100,
    max_depth=15,
    random_state=42,
    class_weight='balanced'
)
rf.fit(X_train_scaled, y_train)
models['Random Forest'] = rf

cv_scores = cross_val_score(rf, X_train_scaled, y_train, cv=5, scoring='roc_auc')

=====
```

```
results['Random Forest'] = {'cv_score': cv_scores.mean()}
print(f"    CV ROC-AUC: {cv_scores.mean():.4f}")
```

3. Training Random Forest...

CV ROC-AUC: 0.7917

```
# 7.4: Gradient Boosting
print("\n4. Training Gradient Boosting...")
gb = GradientBoostingClassifier(
    n_estimators=100,
    learning_rate=0.1,
    max_depth=5,
    random_state=42
)
gb.fit(X_train_scaled, y_train)
models['Gradient Boosting'] = gb

cv_scores = cross_val_score(gb, X_train_scaled, y_train, cv=5, scoring='roc_auc')
results['Gradient Boosting'] = {'cv_score': cv_scores.mean()}
print(f"    CV ROC-AUC: {cv_scores.mean():.4f}")
```

4. Training Gradient Boosting...

CV ROC-AUC: 0.7730

```
# 7.5: XGBoost (Best Model)
print("\n5. Training XGBoost...")
xgb_model = xgb.XGBClassifier(
    learning_rate=0.1,
    max_depth=6,
    n_estimators=200,
    subsample=0.8,
    colsample_bytree=0.8,
    random_state=42,
    eval_metric='logloss'
)
xgb_model.fit(X_train_scaled, y_train)
models['XGBoost'] = xgb_model

cv_scores = cross_val_score(xgb_model, X_train_scaled, y_train, cv=5, scoring='roc_auc')
results['XGBoost'] = {'cv_score': cv_scores.mean()}
print(f"    CV ROC-AUC: {cv_scores.mean():.4f}")
```

5. Training XGBoost...

CV ROC-AUC: 0.8008

```
# Print training summary
print("\n" + "="*50)
```

```
print("TRAINING SUMMARY")
print("*50)

for name, result in results.items():
    print(f"{name}: {result['cv_score']:.4f}")

# Find best model
best_model_name = max(results, key=lambda x: results[x]['cv_score'])
best_model = models[best_model_name]

print("\n✓ Best Model: {best_model_name}")
print("CV ROC-AUC: {results[best_model_name]['cv_score']:.4f}")
```

```
=====
TRAINING SUMMARY
=====
Logistic Regression: 0.8042
Decision Tree: 0.6118
Random Forest: 0.7917
Gradient Boosting: 0.7730
XGBoost: 0.8008

✓ Best Model: Logistic Regression
CV ROC-AUC: 0.8042
```

8. Model Evaluation

```
# =====
# STEP 8: EVALUATE ALL MODELS
# =====

print("*50)
print("MODEL EVALUATION")
print("*50)

evaluation_results = []

for name, model in models.items():
    # Make predictions
    y_pred = model.predict(X_test_scaled)
    y_pred_proba = model.predict_proba(X_test_scaled)[:, 1]

    # Calculate metrics
    metrics = {
        'Model': name,
        'Accuracy': accuracy_score(y_test, y_pred),
        'Precision': precision_score(y_test, y_pred),
        'Recall': recall_score(y_test, y_pred),
```

```

        'F1-Score': f1_score(y_test, y_pred),
        'ROC-AUC': roc_auc_score(y_test, y_pred_proba)
    }

    evaluation_results.append(metrics)

# Create comparison DataFrame
results_df = pd.DataFrame(evaluation_results)

print("\nModel Performance Comparison:")
print(results_df.to_string(index=False))

=====
MODEL EVALUATION
=====

Model Performance Comparison:
      Model  Accuracy   Precision   Recall   F1-Score   ROC-AUC
Logistic Regression  0.748299  0.363636  0.765957  0.493151  0.802998
      Decision Tree  0.785714  0.326087  0.319149  0.322581  0.622104
      Random Forest  0.833333  0.375000  0.063830  0.109091  0.751400
      Gradient Boosting  0.836735  0.476190  0.212766  0.294118  0.750108
      XGBoost  0.857143  0.666667  0.212766  0.322581  0.741752

```

```

# Detailed evaluation of best model
print("\n" + "="*50)
print(f"DETAILED EVALUATION: {best_model_name}")
print("="*50)

y_pred_best = best_model.predict(X_test_scaled)
y_pred_proba_best = best_model.predict_proba(X_test_scaled)[:, 1]

print("\nClassification Report:")
print(classification_report(y_test, y_pred_best,
                           target_names=['No Attrition', 'Attrition']))

=====
DETAILED EVALUATION: Logistic Regression
=====


```

```

Classification Report:
      precision   recall   f1-score   support
No Attrition       0.94     0.74     0.83      247
Attrition         0.36     0.77     0.49       47

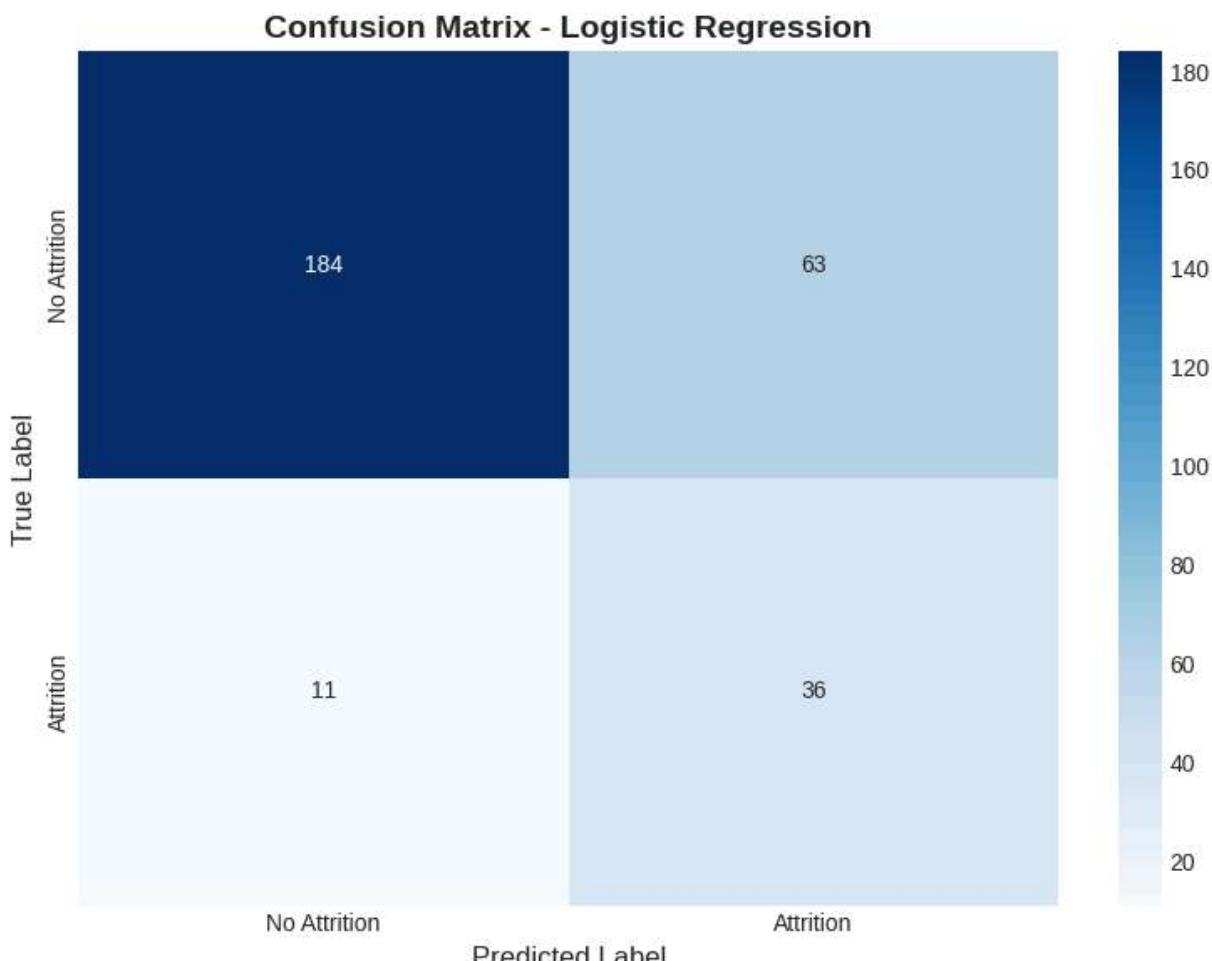
      accuracy           0.75      294
      macro avg          0.65     0.76     0.66      294
      weighted avg        0.85     0.75     0.78      294

```

```
# Confusion Matrix
cm = confusion_matrix(y_test, y_pred_best)

plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['No Attrition', 'Attrition'],
            yticklabels=['No Attrition', 'Attrition'])
plt.title(f'Confusion Matrix - {best_model_name}', fontsize=14, fontweight='bold')
plt.ylabel('True Label', fontsize=12)
plt.xlabel('Predicted Label', fontsize=12)
plt.tight_layout()
plt.show()

print("\nConfusion Matrix:")
print(f" True Negatives: {cm[0,0]}")
print(f" False Positives: {cm[0,1]}")
print(f" False Negatives: {cm[1,0]}")
print(f" True Positives: {cm[1,1]}")
```



▼ 9. SHAP Explainability

```
# =====
# STEP 9: SHAP EXPLAINABILITY
# =====

print("*"*50)
print("SHAP EXPLAINABILITY ANALYSIS")
print("*"*50)
```

```
import shap

# Initialize SHAP explainer for linear models
explainer = shap.LinearExplainer(best_model, X_train_scaled)

# Calculate SHAP values
shap_values = explainer.shap_values(X_test_scaled)

# For binary classification, use values for class 1
if isinstance(shap_values, list):
    shap_values = shap_values[1]

print("✅ SHAP values calculated!")
```

```
=====
SHAP EXPLAINABILITY ANALYSIS
=====
✅ SHAP values calculated!
```

```
# Feature Importance (Mean Absolute SHAP values)
importance = np.abs(shap_values).mean(axis=0)

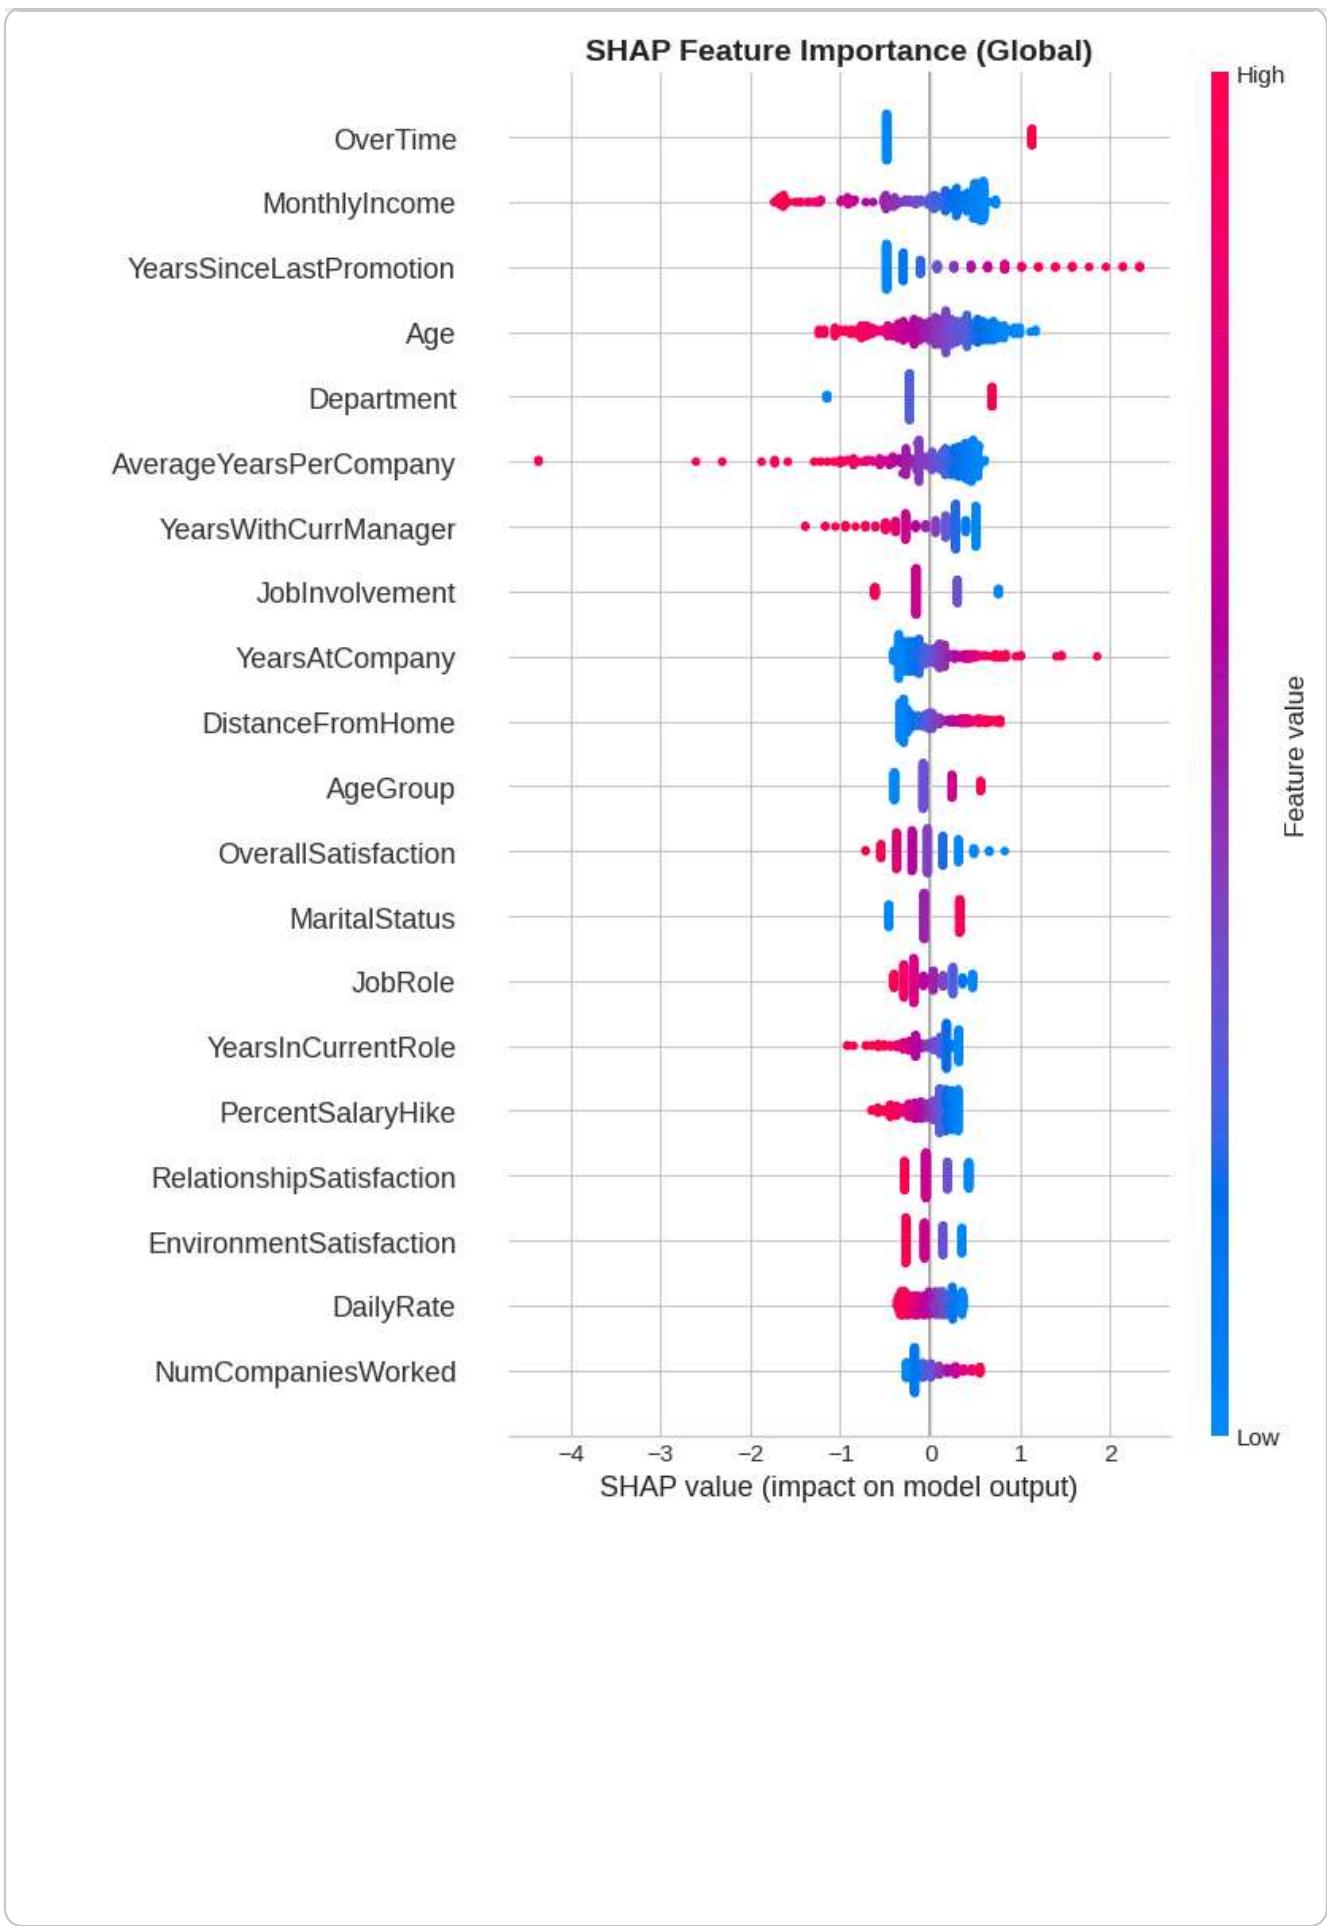
feature_importance = pd.DataFrame({
    'Feature': X.columns,
    'Importance': importance
}).sort_values('Importance', ascending=False)

print("\nTop 10 Most Important Features:")
print(feature_importance.head(10).to_string(index=False))
```

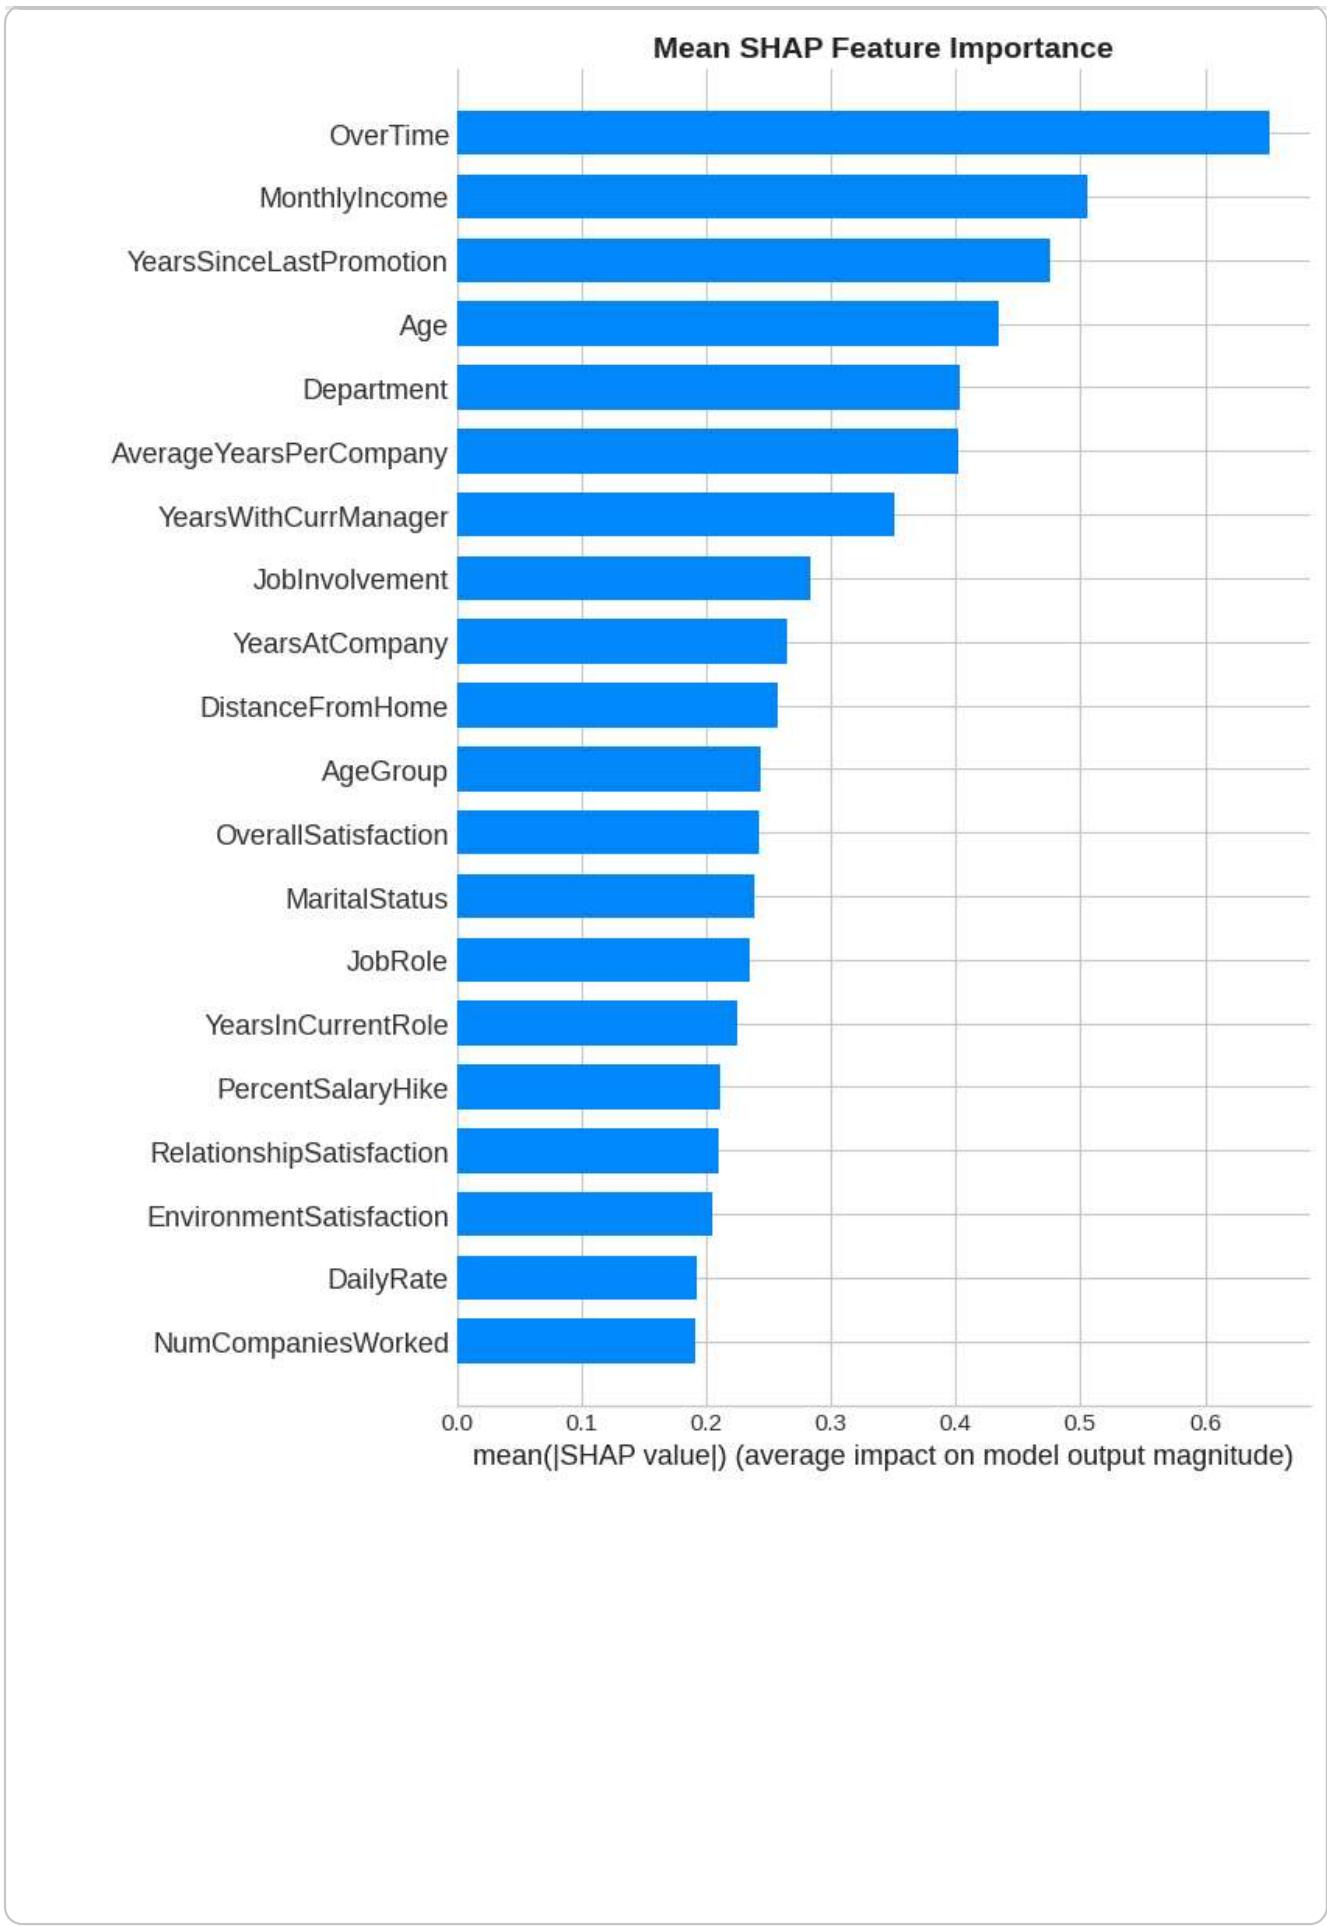
Top 10 Most Important Features:

Feature	Importance
OverTime	0.651626
MonthlyIncome	0.505159
YearsSinceLastPromotion	0.475348
Age	0.434405
Department	0.403781
AverageYearsPerCompany	0.401640
YearsWithCurrManager	0.350982
JobInvolvement	0.283521
YearsAtCompany	0.265264
DistanceFromHome	0.257136

```
# SHAP Summary Plot (Global Feature Importance)
plt.figure(figsize=(10, 8))
shap.summary_plot(shap_values, X_test_scaled, show=False)
plt.title('SHAP Feature Importance (Global)', fontsize=14, fontweight='bold')
plt.tight_layout()
plt.show()
```



```
# SHAP Bar Plot (Mean Absolute Importance)
plt.figure(figsize=(10, 8))
shap.summary_plot(shap_values, X_test_scaled, plot_type="bar", show=False)
plt.title('Mean SHAP Feature Importance', fontsize=14, fontweight='bold')
plt.tight_layout()
plt.show()
```

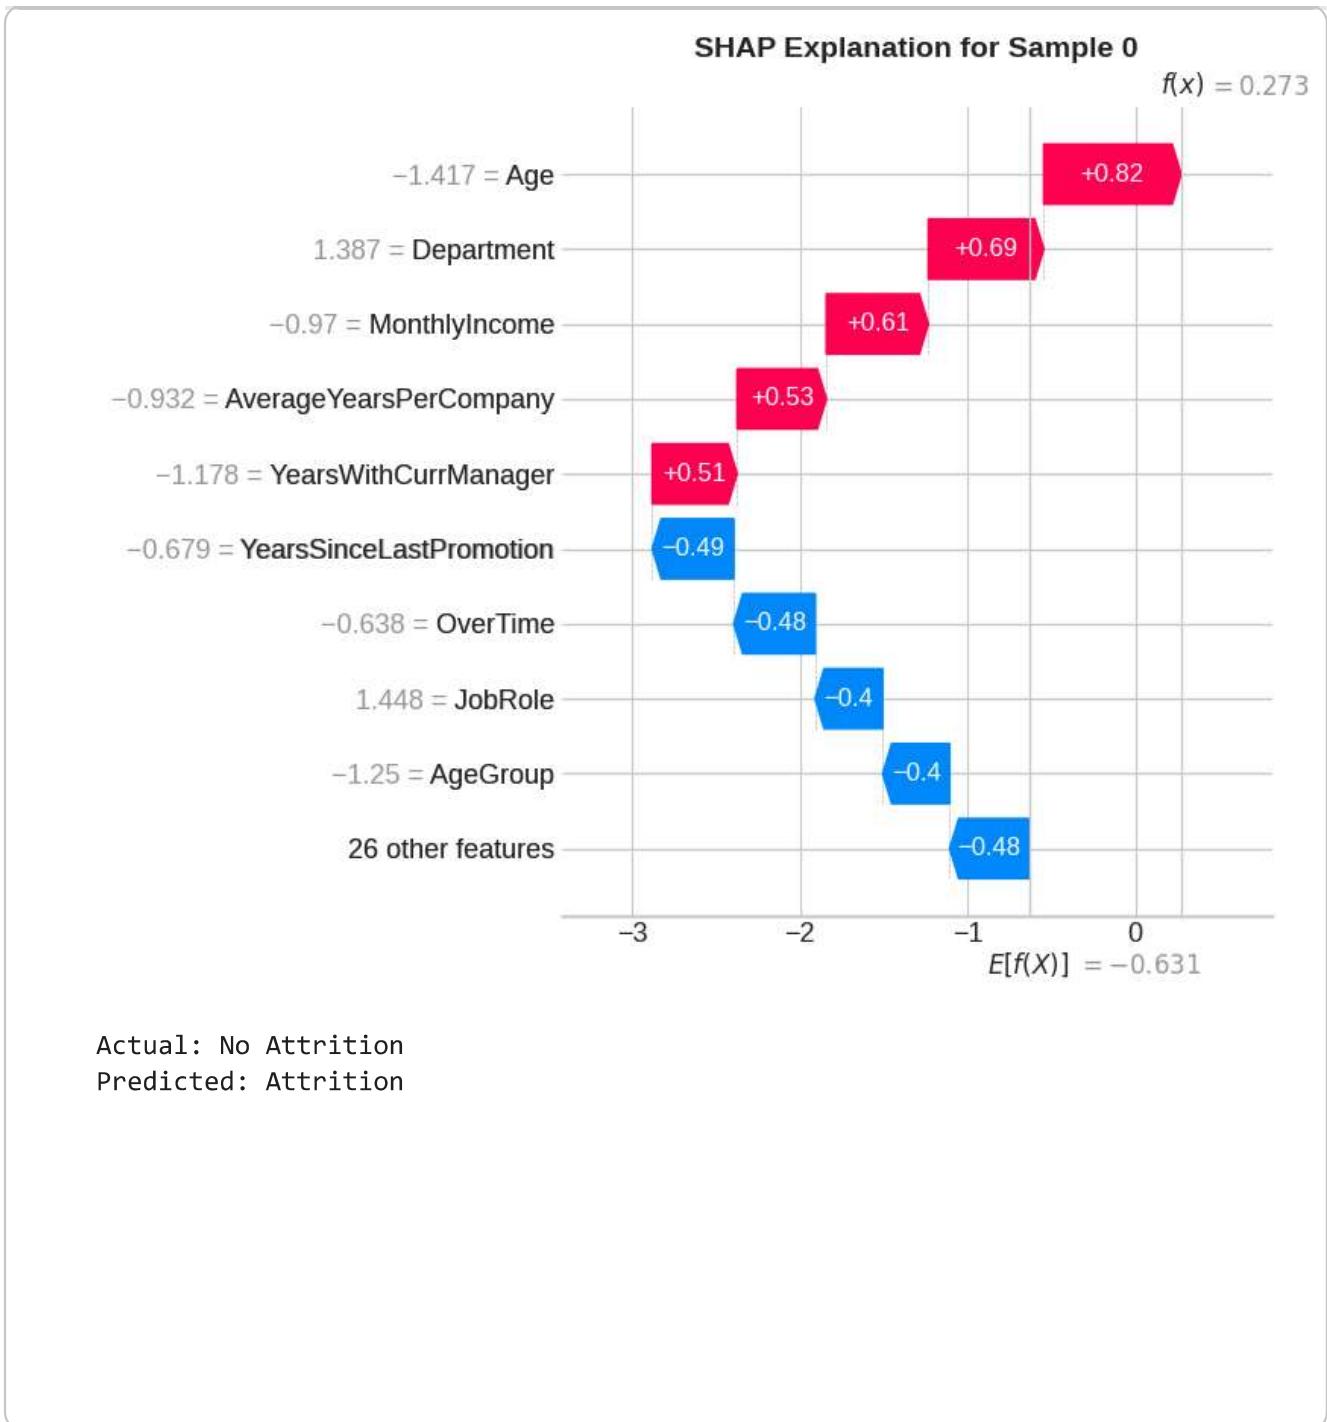


```
# Explain a single prediction (Waterfall Plot)
sample_idx = 0 # Change this to explore different samples

explanation = shap.Explanation(
    values=shap_values[sample_idx],
    base_values=explainer.expected_value,
    data=X_test_scaled.iloc[sample_idx],
    feature_names=X.columns
)

plt.figure(figsize=(12, 8))
shap.waterfall_plot(explanation, show=False)
plt.title(f'SHAP Explanation for Sample {sample_idx}', fontsize=14, fontweight='bold')
plt.tight_layout()
plt.show()

actual = y_test.iloc[sample_idx]
predicted = y_pred_best[sample_idx]
print(f"\nActual: {'Attrition' if actual == 1 else 'No Attrition'}")
print(f"Predicted: {'Attrition' if predicted == 1 else 'No Attrition'}")
```



Actual: No Attrition

Predicted: Attrition

▼ 10. Visualization

```
# =====
# STEP 10: CREATE VISUALIZATIONS
# =====

print("*50)
print("CREATING VISUALIZATIONS")
print("*50)
```

```
# 10.1: Model Performance Comparison
fig, axes = plt.subplots(2, 3, figsize=(15, 10))
axes = axes.flatten()

metrics = ['Accuracy', 'Precision', 'Recall', 'F1-Score', 'ROC-AUC']

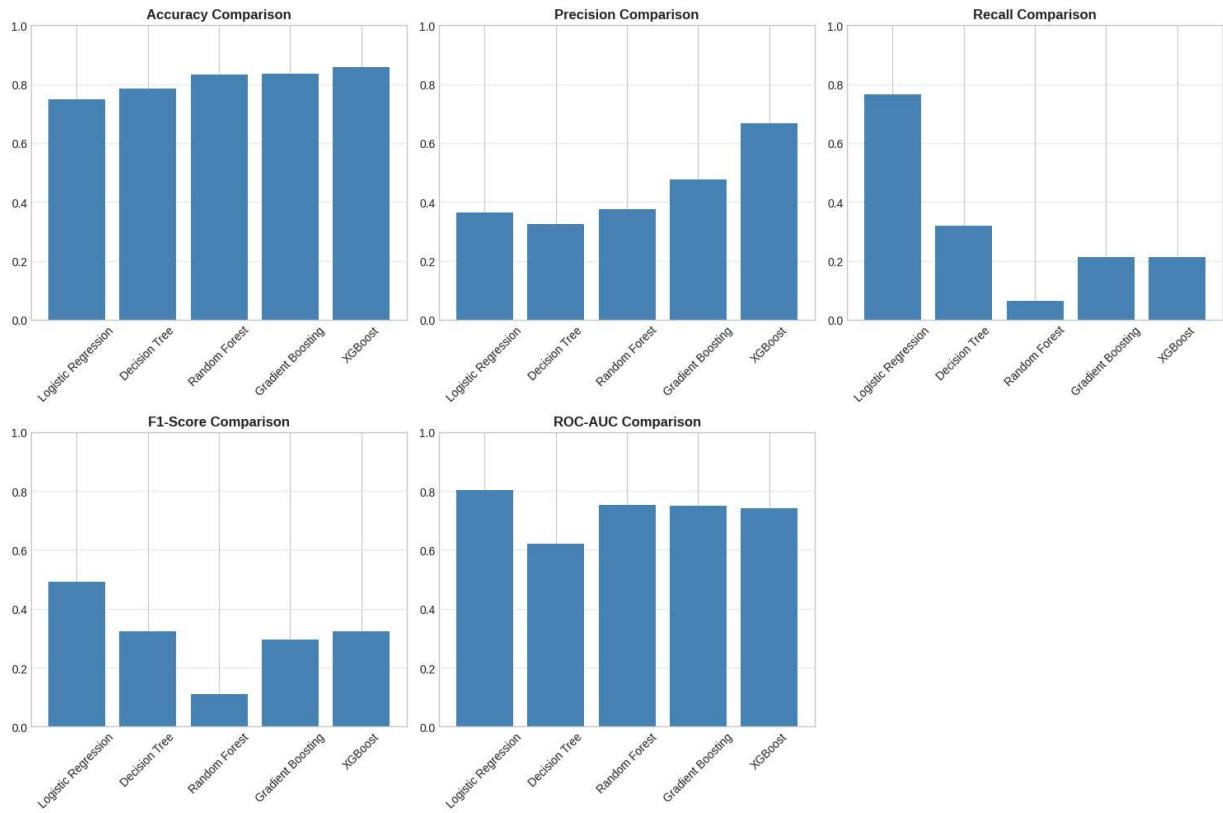
for i, metric in enumerate(metrics):
    axes[i].bar(results_df['Model'], results_df[metric], color='steelblue')
    axes[i].set_title(f'{metric} Comparison', fontsize=12, fontweight='bold')
    axes[i].set_ylim([0, 1])
    axes[i].tick_params(axis='x', rotation=45)
    axes[i].grid(axis='y', alpha=0.3)

axes[5].axis('off')
plt.tight_layout()
plt.show()
```

=====

CREATING VISUALIZATIONS

=====

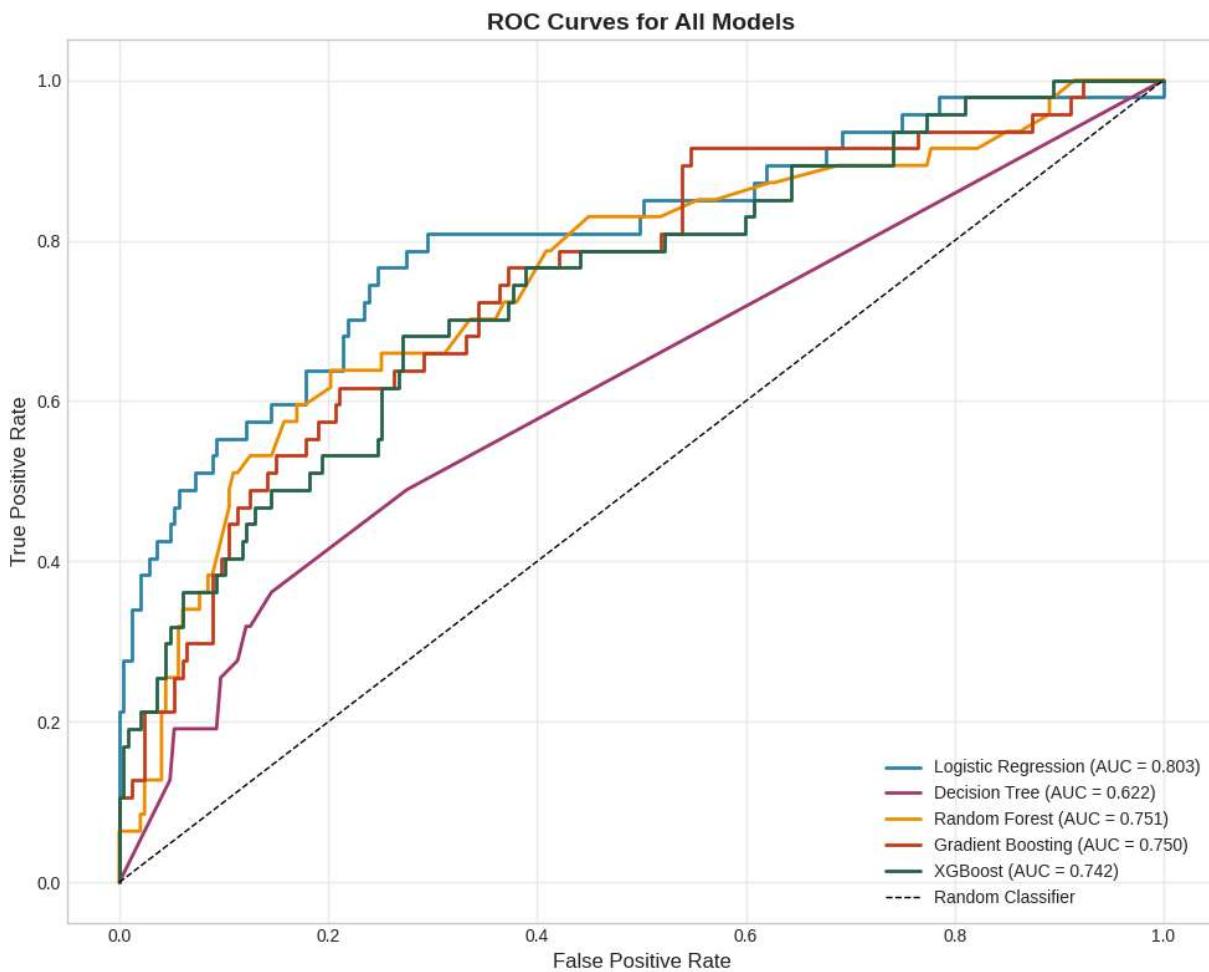


```
# 10.2: ROC Curves for All Models
plt.figure(figsize=(10, 8))

colors = ['#2E86AB', '#A23B72', '#F18F01', '#C73E1D', '#276749']

for i, (name, model) in enumerate(models.items()):
    y_pred_proba = model.predict_proba(X_test_scaled)[:, 1]
    fpr, tpr, _ = roc_curve(y_test, y_pred_proba)
    auc = roc_auc_score(y_test, y_pred_proba)
    plt.plot(fpr, tpr, linewidth=2, label=f'{name} (AUC = {auc:.3f})', color=colors[i])

plt.plot([0, 1], [0, 1], 'k--', linewidth=1, label='Random Classifier')
plt.xlabel('False Positive Rate', fontsize=12)
plt.ylabel('True Positive Rate', fontsize=12)
plt.title('ROC Curves for All Models', fontsize=14, fontweight='bold')
plt.legend(loc='lower right')
plt.grid(alpha=0.3)
plt.tight_layout()
plt.show()
```

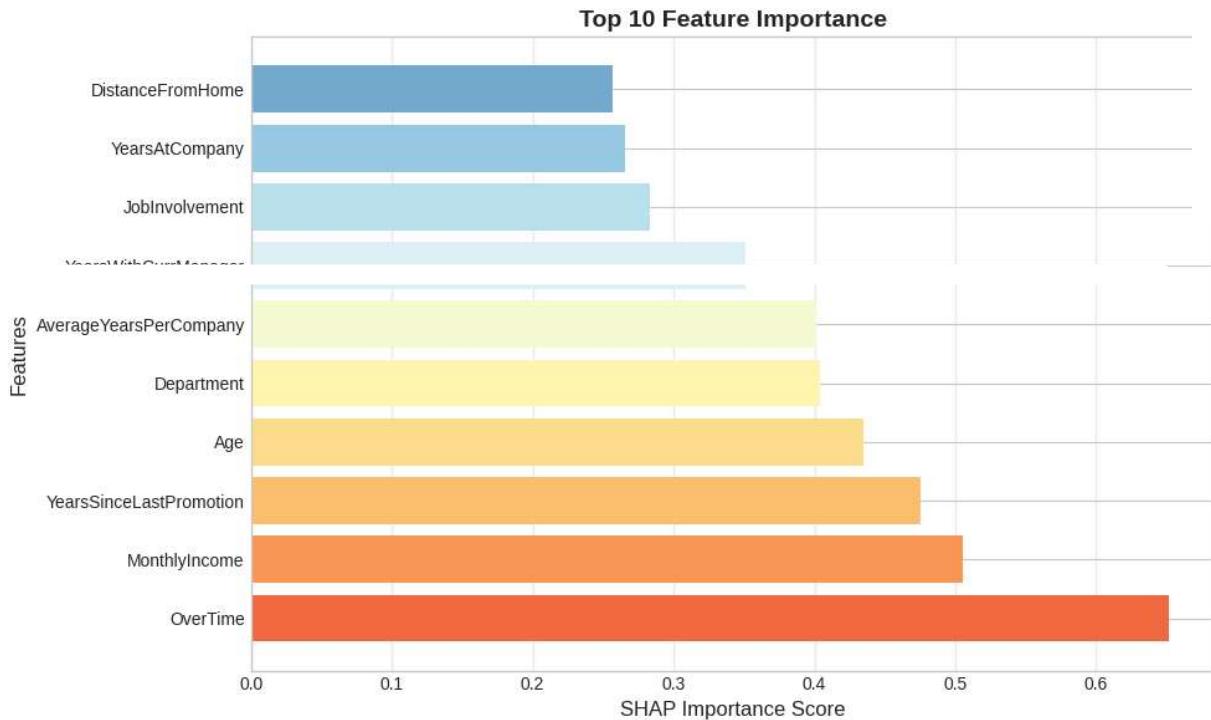


```
# 10.3: Feature Importance Bar Chart
plt.figure(figsize=(10, 6))

top_features = feature_importance.head(10)
colors = plt.cm.RdYlBu(np.linspace(0.2, 0.8, len(top_features)))

plt.barh(top_features['Feature'], top_features['Importance'], color=colors)
```

```
plt.xlabel('SHAP Importance Score', fontsize=12)
plt.ylabel('Features', fontsize=12)
plt.title('Top 10 Feature Importance', fontsize=14, fontweight='bold')
plt.grid(axis='x', alpha=0.3)
plt.tight_layout()
```



▼ 11. Save Model

```
# =====
# STEP 11: SAVE MODEL AND PREPROCESSOR
# =====
```